CHAIR OF DECENTRALIZED INFORMATION SYSTEMS & DATA MANAGEMENT

TECHNICAL UNIVERSITY OF MUNICH

Bachelor's Thesis in Informatics

Concurrent Range Locking

Thua-Duc Nguyen

CHAIR OF DECENTRALIZED INFORMATION SYSTEMS & DATA MANAGEMENT

TECHNICAL UNIVERSITY OF MUNICH

Bachelor's Thesis in Informatics

Concurrent Range Locking

Author:	Thua-Duc Nguyen
Supervisor:	Prof. Dr. Viktor Leis
Advisor:	Lam-Duy Nguyen
SubmissionDate:	15.09.2024

I confirm that this bachelor's thesis in informatics is my own work and I have documented all sources and material used.

München, 15.09.2024   Thua-Duc Nguyen

Acknowledgments

Abstract

In modern computing environments, managing concurrent access to shared resources, such as database tables, file regions, or memory segments, is a significant challenge. Traditional synchronization methods often rely on a single lock, leading to inefficiencies and performance bottlenecks under high concurrency. Range locking offers a more refined solution by partitioning shared resources into multiple segments, allowing different processes to access these segments concurrently.

Although range locking plays a crucial role in various systems, including database management systems (DBMS), file systems, and operating systems, it remains an under-studied topic. Existing implementations often suffer from contention and inefficiencies, particularly under high concurrency, highlighting the need for more scalable solutions.

This thesis introduces a new lock-free concurrent range-locking mechanism to overcome these challenges. Our method improves upon previous designs by eliminating bottlenecks and ensuring high performance in heavily concurrent environments. Our evaluation shows that the proposed method outperforms existing approaches. This thesis provides a comprehensive assessment of the development and evaluation of this new range lock, offering a comprehensive assessment of its effectiveness in modern computing systems.

Contents

List of Figures

List of Tables

1 Introduction

In modern computing environments, managing concurrent access to shared resources, such as database tables, file regions, or memory segments, is a significant challenge. Systems often rely on coarser synchronization methods, like using a single lock to manage the access between threads. This leads to significant inefficiencies and performance bottlenecks at high concurrency levels. Range lock [1, 2, 3] offers a more refined solution to this challenge. It partitions a shared resource into multiple arbitrarily sized segments, allowing different processes to exclusively acquire and access each segment simultaneously. This strategy effectively addresses the drawbacks and bottlenecks associated with single-lock methods, provides more granular control, and improves parallelism.

Range locking is essential across various systems. Despite being an under-researched topic, range locking plays a critical role in several key areas, including database management systems (DBMS), file systems, and operating systems.

In DBMS, range locks ensure data consistency and prevent anomalies such as phantoms, especially at the serializable isolation level [4]. They secure not only individual records, but also the gaps between them, preventing other transactions from inserting or modifying records within the range until the transaction is complete. This is particularly challenging in systems where transaction control (TC) and data control (DC) are separated. TC must lock the range before safely interacting with DC, despite not knowing the specific keys involved.

Similarly, file systems, especially those designed for high-performance computing, are required to handle large-scale parallel I/O requests targeting different parts of a single data file [5, 6]. To achieve this, these systems must efficiently and precisely coordinate concurrent access to shared storage, ensuring consistency and preventing conflicts.

Within operating systems, particularly in the Linux kernel community, increasing interest has been in applying range locks to alleviate contention issues related to `mmap_sem`. `mmap_sem` is a semaphore used to manage access to virtual memory areas (VMAs) [7, 8, 9]. Because VMAs require synchronized access for operations such as

memory mapping, unmapping, and handling page faults, range locking could come into play.

Existing range lock is not efficient. Previous implementations of range-locking mechanisms often are not efficient. These implementations often suffer from contention points due to the reliance on a single lock [10, 3], as spinlock effectively serializes all incoming lock and unlock requests. Additionally, some methods may be complex and tightly coupled with lock-based concurrency control protocols, which are not applicable for general DBMS operations [11, 12]. These limitations highlight the necessity for more refined and scalable solutions to better accommodate the demands of modern, large-scale systems. This drives us to develop an enhanced range lock.

Contribution. In the scope of this thesis, we introduce a new lock-free concurrent range lock. Our method addresses the inefficiencies of previous range locks and ensures high performance in highly concurrent environments. The result shows that our method is at least three times faster than existing solutions. This thesis provides an in-depth exploration of the development and evaluation of the proposed range-locking mechanism, offering a thorough understanding of its performance. Furthermore, we compare the proposed solution with state-of-the-art methods, clearly assessing its effectiveness.

2 Related Work

2.1 Coarse-Grained Range Lock

2.1.1 Tree-Based Range Lock

Several works have explored coarse-grained range-locking methods. Jan Kara introduced a range-locking mechanism for the Linux kernel [10], which utilizes a range tree (specifically a red-black tree) to manage range locks and employs a spinlock for synchronization. Each lock is represented as a node in the tree. Similarly, Kim et al. adopted a comparable range lock in their work on pNOVA [13], a variant of the NOVA file system that uses range-based reader-writer locks to enable parallel I/O within a single shared file.

When a thread requests a range lock, it first acquires a spinlock, then traverses the tree to determine the number of locks intersecting with the requested range. Afterward, the thread inserts a node describing its range into the tree and releases the spinlock. If no intersecting locks are found, the thread can proceed with accessing the critical section. If intersecting locks are detected, the thread waits until those locks are released and the number of intersecting locks drops to zero. Upon completing its operation, the thread re-acquires the spinlock, removes its node from the tree, updates the count of overlapping locks, and releases the spinlock. This method ensures that each range is locked only after all previous conflicting range locks have been released, thereby achieving fairness and avoiding livelocks.

Drawbacks

One significant observation is that the coarse-grained spinlock of an interval tree can severely hinder parallelism, as the spinlock effectively serializes all incoming lock and unlock requests. Under heavy concurrent access, this serialization easily becomes a contention point, limiting the system's performance.

Consider three exclusive lock requests for the ranges $A = [1..3]$, $B = [2..7]$, and $C = [4..5]$, arriving in that order. While A holds the lock, B is blocked because it overlaps

with A, and C is blocked behind B. However, in practice, C does not overlap with A and could proceed without waiting. This unnecessary blocking reduces the overall efficiency and concurrency of the system.

2.1.2 List-Based Range Lock

Song et al. [3] introduced a dynamic range-locking design to enhance the implementation of the Linux kernel. Their range lock uses a skip list [14] to dynamically manage the address ranges that are currently locked.

When a thread requests a specific range `[start, start+len)`, the range lock searches the skip list. If an existing or overlapping range is found, it indicates that another thread is currently modifying that range, requiring the requesting thread to wait and then retry. If no overlapping range is found, the requested range is added to the skip list, signifying that the lock has been acquired. Releasing a range involves deleting the corresponding range from the skip list.

Compared to the interval tree, the skip list is more lightweight and efficient, allowing for quicker searches of overlapping ranges.

Drawbacks

Similar to the tree-based range lock, contention remains an issue with this approach. Additionally, it unnecessarily blocks non-conflicting requests, further reducing system efficiency and limiting concurrency.

2.2 Fine-Grained Range Lock

2.2.1 List-Based Range Lock

Kogan et al. [2] introduced a range lock based on a concurrent linked list, where each node represents an acquired range. This design aims to provide a lock-free method, addressing critical shortcomings of previous range-locking implementations. In a lock-free system, processes can proceed without being blocked by locks held by other processes, thereby improving performance and scalability.

The proposed method involves inserting acquired ranges into a linked list sorted by their starting points, ensuring that only one range from a group of overlapping ranges can be inserted using an atomic `compare-and-swap` (CAS) operation. A significant

difference in this method compared to previous ones is that each node has two statuses: `marked` (logically deleted) or `unmarked` (present).

When a thread wants to acquire a range, it iterates through the skip list. If it encounters a marked node, it removes it using CAS and continues to iterate. If the current node protects a range that overlaps, the thread waits until that node is deleted. Otherwise, a node is inserted into the list, signaling that the range is acquired. To release a range, the thread marks the node as deleted.

Drawbacks

Linked List Inefficiency: While this design implements a lock-free method that effectively addresses the limitations of existing range locks, it comes with its own set of trade-offs. In general, insertion and lookup operations in a linked list are less efficient than in tree-like structures. The average time complexity for searching in a linked list is $O(n)$, whereas it is only $O(\log n)$ for skip lists or tree-like structures [15]. Our evaluation will demonstrate that this inefficiency becomes particularly pronounced when handling multiple overlapping ranges within the list.

2.2.2 Bitmap Range Lock

In addition to the tree-based range locking method discussed in Subsection 2.1.1, Kim et al. proposed a lock-free range lock, which they claim offers enhanced efficiency compared to interval tree-based locks. This approach involves dividing a file into segments, each managed by a 32-bit variable that functions as a reader-writer lock. The most significant bit represents the writer lock status (1 for locked, 0 for unlocked), while the remaining 31 bits count the number of active readers. The method utilizes hardware-supported atomic operations to ensure that writer locks can only be set when no other locks are active and that reader locks are granted as long as no writer lock is present. Unlocking is achieved by clearing the writer lock bit and decrementing the reader counter.

Although this method provides finer-grained locking with reduced overhead compared to interval tree-based locks, it is tailored to handle both reader and writer modes and depends on specific memory size constraints. Since our research focuses exclusively on the exclusive mode and does not address the reader-writer combination, we have chosen not to consider this approach as a competitor in our project.

3 Data structure

3.1 Skiplist is suitable for range lock

The effectiveness of a range lock is strongly influenced by the underlying data structure used to manage these ranges. Existing range locks have used many different comparison-based data structures, such as interval tree [10, 13] and linked list [3], skip list [2] each of which has advantages and disadvantages.

Interval trees are considered heavyweight compared to linked lists and skip lists because they require complex balancing mechanisms to maintain efficient search times, which adds significant computational and memory overhead. Interval trees also require the maintenance of additional data, such as subtree intervals, making them more resource-intensive. This makes interval trees unsuitable for range locking.

Linked lists have minimal structure and are easy to implement and maintain. However, it is very expensive in terms of Insert/Delete time. The complexity O(n) for the average case of search becomes a significant disadvantage when the number of ranges increases,

Skiplists strike a balance between the heaviness of interval trees and the simplicity of linked lists. Despite their theoretically poor worst-case performance, skip lists rarely exhibit worst-case behavior, making them efficient in most scenarios. For example, in a dictionary with more than 250 entries, the probability that a search will take more than three times the expected time is less than one in a million [16].

We strongly believe that the skip list is ideal for implementing range locking.

Data Structure	Best Case	Average Case	Worst Case
Linked List	$O(1)$	$O(n)$	$O(n)$
Skip List	$O(1)$	$O(logn)$	$O(n)$
Interval Tree	$O(\log n)$	$O(\log n)$	$O(\log n)$

Table 3.1: Time complexities of Insert/Delete operations in Linked List, Skip List and Interval Tree

Figure 3.1: An example of a skip list. It has five levels of sorted linked lists.

3.2 Modification of skiplist for concurrent range lock

We developed our concurrent range lock based on the `LockFreeSkipList` proposed by Herlihy et al. [17]. In summary, `LockFreeSkipList` uses atomic operations (`compareAndSet()`) to manage `Node` references without locks, which enhances performance in multithreaded environments. When adding a `Node`, the process starts at the lowest level and moves upward to ensure immediate visibility. Removing a `Node` involves marking nodes from the top down before unlinking them. Furthermore, it relaxes strict structural maintenance of higher levels, focusing on the bottom-level list for set representation, which offers improved scalability and efficiency.

However, conventional skiplists does not work out of the box for range locking. Firstly, each Node of the skiplist stores key and value, not range. Secondly, the comparison condition needs to be modified because we now must compare both range start and range end while inserting or deleting instead of just using the key. We present this in two critical changes in the Sections 3.2.1 and 4.3.3, respectively.

3.2.1 Node

Node is the base of our `ConcurrentRangeLock` structure. Each `Node` now contains `start` and `end`, which represents range. `Node` uses an array of `AtomicMarkableReference` (more details in Section 3.2.2) to maintain forward links at each level, which allows for efficient traversal and updates. `Node` provides the following methods:

- `initialize`: sets up a `Node` with specific range and level values.

- `initializeHead`: configures the head `Node` with forward pointers directed to a provided tail `Node`, establishing the initial structure.

- `getTopLevel()`, `getStart()`, `getEnd()`: accessor methods to retrieve the Node's properties.

```cpp
class Node {
  private:
    uint64_t start, end;
    int topLevel;
    AtomicMarkableReference** next = nullptr;

  public:
    void initialize(uint64_t start, uint64_t end, int topLevel) {
        this->start = start; this->end = end;
        this->topLevel = topLevel;

        next = new AtomicMarkableReference*[topLevel + 1];
        for (int i = 0; i <= topLevel; ++i) {
            next[i] = new AtomicMarkableReference();
        }
    }

    void initializeHead(uint64_t start, uint64_t end, int topLevel, Node* tail) {
        initialize(start, end, topLevel);
        for (int i = 0; i <= topLevel; ++i){
            next[i]->store(tail, false);
        }
    }

    int getTopLevel() const { return topLevel; }
    uint64_t getStart() const { return start; }
    uint64_t getEnd() const { return end; }
};
```

Listing 3.1: Pseudocode for Node structure

3.2.2 AtomicMarkableReference

The `AtomicMarkableReference` class uses a single atomic variable, `atomicRefMark`, to store a packed representation of both the reference (specifically, a `Node`) and a mark. If the mark is 1, it indicates that the `Node` it references is softly deleted. These values are packed and unpacked using bitwise operations, where the least significant bit represents the mark.

Listing 3.2 provides the pseudo code for `AtomicMarkableReference`.

The `pack` method combines a `Node` pointer and a boolean mark into a single `uintptr_t` value by encoding the pointer into the lower bits and the mark into the highest bit. Conversely, the `unpack` method decodes this packed value to retrieve the original `Node` pointer and boolean mark.

To atomically set a new `Node` pointer and mark value, the `store` method uses relaxed memory ordering.

The `compareAndSet` method performs an atomic update of both the reference and mark if they match the expected values, employing acquire-release memory ordering for proper synchronization.

The `attemptMark` method focuses on updating the mark alone, provided that the current reference matches the expected one and the mark differs. If the update succeeds, it returns `true`; otherwise, it returns `false`.

The `get` method retrieves the current reference and mark, which stores the mark in a provided boolean pointer. In contrast, the `getReference` method returns the current reference without accessing the mark.

```cpp
class AtomicMarkableReference {
  private:
    std::atomic<uintptr_t> atomicRefMark;

    uintptr_t pack(Node* ref, bool mark) const {
        return reinterpret_cast<uintptr_t>(ref) | (mark ? 1 : 0);
    }

    std::pair<Node*, bool> unpack(uintptr_t packed) const {
        return {reinterpret_cast<Node*>(packed & -1), packed & 1};
    }

  public:
    AtomicMarkableReference() {
        atomicRefMark.store(pack(nullptr, false), std::memory_order_relaxed);
    }

    void store(Node* ref, bool mark) {
        atomicRefMark.store(pack(ref, mark), std::memory_order_relaxed);
    }

    bool compareAndSet(Node* expectedRef, Node* newRef, bool expectedMark, bool newMark) {
        return atomicRefMark.compare_exchange_strong(
            pack(expectedRef, expectedMark), pack(newRef, newMark), std::memory_order_acq_rel);
    }

    bool attemptMark(Node* expectedRef, bool newMark) {
        auto [currentRef, currentMark] = unpack(atomicRefMark.load(std::memory_order_acquire));
        if (currentRef == expectedRef && currentMark != newMark) {
            return atomicRefMark.compare_exchange_strong(
                current, pack(expectedRef, newMark), std::memory_order_acq_rel);
        }
        return false;
    }

    Node* get(bool* mark) const {
        auto [ref, currentMark] = unpack(atomicRefMark.load(std::memory_order_acquire));
        mark[0] = currentMark;
        return ref;
    }

    Node* getReference() const {
        auto [ref, _] = unpack(atomicRefMark.load(std::memory_order_acquire));
        return ref;
    }
};
```

Listing 3.2: AtomicMarkableReference

4 Approach

In this chapter, we will provide a detailed explanation of the API and algorithm behind the concurrent range lock.

4.1 Concurrent Range Lock API

Figure 4.1: ConcurrentRangeLock. Each `Node` has a unique range.

The ConcurrentRangeLock class consists of two main functions: `tryLock` and `releaseLock`.

The `tryLock` method attempts to acquire a lock for the specified range [`start`, `end`], returning `true` on success and `false` otherwise. The `releaseLock` method releases the lock for the range [`start`, `end`], with `true` indicating success and `false` if the range was not found or an error occurred. We will discuss these methods in Subsection 4.3.1 and 4.3.2.

The two primary methods rely heavily on private searching methods such as `findInsert`, `findExact`, and `findDelete`, which handle insertion finding, exact range finding, and physical deletion of ranges, respectively. We will discuss these methods in Section 4.3.3.

```
1  class ConcurrentRangeLock {
2    public:
3      ConcurrentRangeLock();
4
5      bool tryLock(uint64_t start, uint64_t end);
6      bool releaseLock(uint64_t start, uint64_t end);
7
8    private:
9      Node *head, *tail;
10
11     int randomLevel();
12     bool findInsert(uint64_t start, uint64_t end, Node **preds, Node **succs);
13     bool findExact(uint64_t start, uint64_t end, Node **preds, Node **succs);
14     void findDelete(uint64_t start, uint64_t end);
15 };
```

Listing 4.1: Pseudocode for ConcurrentRangeLock API

4.2 Life cycle

Each thread has to call the `tryLock` method with the parameters range start, range end, thus the positions on the shared object that this thread wants to acquire. If there is no conflict, that is, if there is no overlap existing on the lock, `tryLock` is successful. Thread now has exclusive access to this segment of the shared object and can freely read and write this segment. After finish using this segment, thread can call `releaseLock` to return his access to this segment. If used correctly, `releaseLock` should always return true.

Figure 4.2: Life cycle of range lock

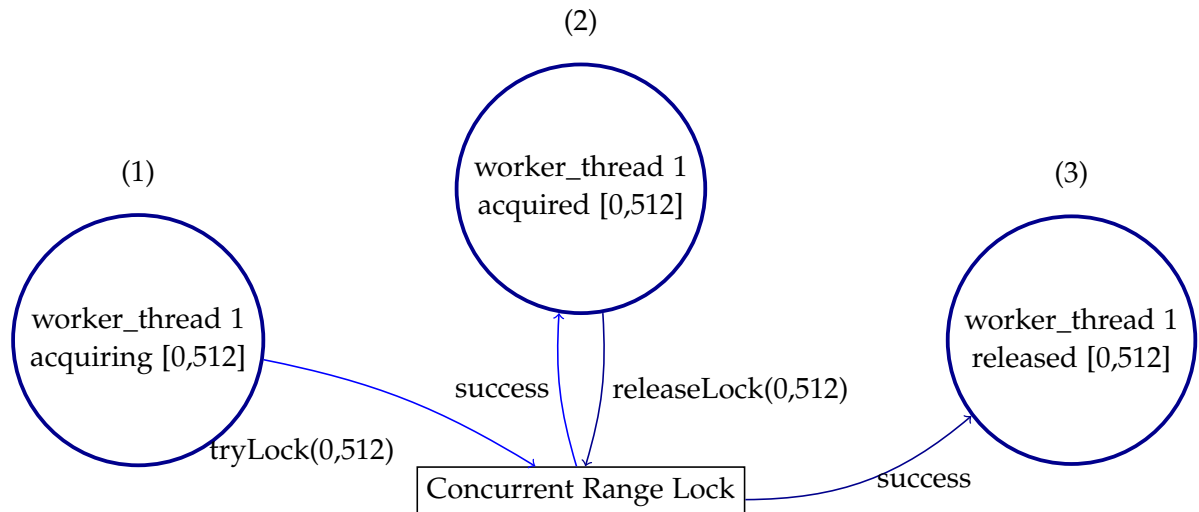## 4.3 Algorithm in details

For the sake of simplicity, we use `uint64_t` for our pseudocode provided in this section. We use the template feature in our open-source C++ implementation to enable generic programming. For further detail, please refer to our open-source code.
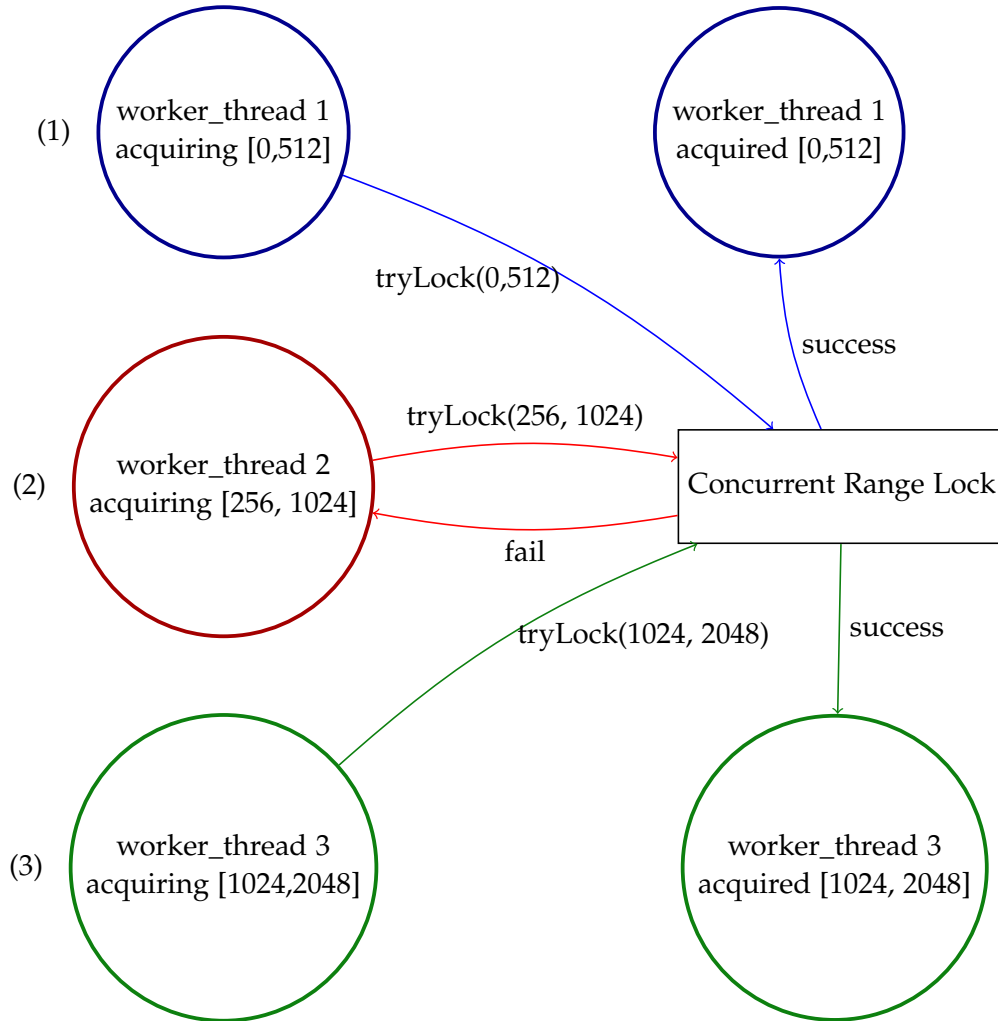

Figure 4.3: The behavior of range lock case: three threads trying to lock ranges. Second thread failed because of overlap

4.3.1 Try Lock

The `tryLock` method, shown in Listing 4.2, utilizes `findInsert()` to check if a node with the range [`start, end`] already exists. If found, `tryLock` returns `false`. Otherwise, it creates a new node and attempts to insert it into the list. The Node is inserted starting from the bottom level, with `compareAndSet()` ensuring the integrity of the insertion. If any insertion fails due to concurrent changes, `findInsert()` is called again to update the predecessors and successors, and the process repeats until successful.

```cpp
bool ConcurrentRangeLock::tryLock(uint64_t start, uint64_t end) {
    int topLevel = randomLevel();
    Node *preds[maxLevel + 1];
    Node *succs[maxLevel + 1];

    while (true) {
        if (findInsert(start, end, preds, succs)) {
            return false;
        } else {
            auto newNode = new Node();
            newNode->initialize(start, end, topLevel);

            for (int level = 0; level <= topLevel; ++level) {
                newNode->next[level]->store(succs[level], false);
            }

            if (!preds[0]->next[0]->compareAndSet(succs[0], newNode, false, false)) {
                continue;
            }

            for (int level = 1; level <= topLevel; ++level) {
                while (!preds[level]->next[level]->compareAndSet(
                    succs[level], newNode, false, false)) {
                    findInsert(start, end, preds, succs);
                }
            }

            return true;
        }
    }
}
```

Listing 4.2: Pseudocode for tryLock method

Algorithm in details

The `tryLock` method, shown in Listing 4.2, uses `findInsert()` to determine whether a node with range [`start, end`] is already in the list (line 7). `tryLock` also calls `findInsert()` to initialize the preds[] and succs[] arrays to hold the new Node's

ostensible predecessors and successors. If an unmarked node is found in the bottom-level list, `findInsert()` returns `true` and the `tryLock` method returns `false`, indicating that part of the target range is already acquired by other thread. See Figure 4.3 as an example. The unsuccessful linearization point of `tryLock` is the same as the successful linearization point of `findInsert()` (line 8). If no node is found, the next step is to add a new node with the key into the structure.

A new node is created with a randomly chosen top-level (lines 10–11). The Node's next references are unmarked and set to the successors returned by the `findInsert()` method (lines 13–15). The next step is to try to add the new Node by linking it into the bottom-level list between the `preds[0]` and `succs[0]` nodes returned by `findInsert()`. We use the `compareAndSet()` method to set the reference while validating that these nodes still refer one to the other and have not been removed from the list (line 17). If the `compareAndSet()` fails, something has changed and the call restarts. If the `compareAndSet()` succeeds, the item is added, and line 17 is the call's linearization point. The findInsert() then links the Node in at higher levels (lines 21–25). For each level, it attempts to splice the Node by setting the predecessor, if it refers to the valid successor, to the new Node (lines 22–23). If successful, it breaks and moves on to the next level. If unsuccessful, then the Node referenced by the predecessor must have changed, and `findInsert()` is called again to find a new valid set of predecessors and successors (line 24). We discard the result of calling `findInsert()` because we care only about recomputing the ostensible predecessors and successors on the remaining unlinked levels. The method returns true once all levels are linked (line 27).

4.3.2 Release Lock

The `releaseLock` method first calls `findExact()` to locate an unmarked node with the specified range [`start, end`]. If found, it marks all node levels except the bottom one, preparing the Node for removal. The method then attempts to mark the bottom-level link using `compareAndSet()`, which serves as the linearization point for a successful `releaseLock`. If marking fails due to concurrent modifications, the method retries or returns false, depending on the Node's state.

Algorithm in details

The `releaseLock` method, shown in Listing 4.3, calls `findExact()` to determine whether an unmarked node with a matching range [`start, end`] is in the bottom-level list (line 7). If no node is found in the bottom-level list, or the Node with a matching

range [start, end] is marked, the method returns false. The linearization point of the unsuccessful releaseLock is that of the findExact() method called in line 7.

If an unmarked node is found, the method logically removes the associated key from the abstract set and prepares it for physical removal. This step uses the set of ostensible predecessors (stored by findExact() in preds[]) and the victim (returned from findExact() in succs[]). First, starting from the top-level, all links up to and **not including** the bottom-level link are marked (lines 12–20) by repeatedly reading next and its mark and applying attemptMark(). If the link is found to be marked (either because it was already marked or because the attempt succeeded), the method moves on to the next-level link. Otherwise, the current level's link is reread since another concurrent thread must have changed it, so the marking attempt must be repeated.

Once all levels but the bottom one have been marked, the method marks the bottom level's next reference. If successful, this marking (line 27) is the linearization point of a successful releaseLock. The releaseLock method tries to mark the next field using compareAndSet(). If successful, it can determine that it was the thread that changed the mark from false to true. Before returning true, the findDelete() method is called. This call is an optimization: findDelete() physically removes all links to the Node it is searching for.

On the other hand, if the compareAndSet() call fails, but the next reference is marked, then another thread must have concurrently removed it, so releaseLock returns false. The linearization point of this unsuccessful releaseLock is the linearization point of the releaseLock method by the thread that successfully marked the next field. Notice that this linearization point must occur during the releaseLock call because the findExact() call found the Node unmarked before it found it marked.

Finally, if the compareAndSet() fails and the Node is unmarked, the next Node must have changed concurrently. Since the victim is known, there is no need to call find() again, and releaseLock simply uses the new value read from next to retry the marking.

4.3.3 Find

Both tryLock and releaseLock methods rely heavily on find methods. There are several find methods in our implementation that serve different purposes:

- bool findInsert(uint64_t start, uint64_t end, Node** preds, Node** succs): checks if the target range [start, end] is free to be inserted.

- `bool findExact(uint64_t start, uint64_t end, Node** preds, Node** succs)`: checks if the target range [`start, end`] is already present in the skip list.

- `void findDelete(uint64_t start, uint64_t end)`: finds the target range [`start, end`] from the skip list to physically delete the Node which contains the corresponding range.

These `findInsert` and `findExact` methods also fill in the `preds[]` and `succs[]` arrays with the target node's predecessors and successors at each level. Because the goal of `findDelete` is only to snip out all the deleted Node, there is no need to fill any array.

Nevertheless, these methods have to maintain the following two properties:

- During traversal, they need to skip over marked nodes. They use `compareAndSet()` to ensure that they remove all softly deleted Node on the way.

- Every `preds[]` reference is to a node with a key strictly less than the target.

Algorithm in details

The `find()` method starts by traversing the LockFreeSkipList from the `topLevel` of the head sentinel, which has the maximal allowed node level. It proceeds down the list level by level, filling in the `preds` and `succs` nodes. These nodes are repeatedly advanced until `pred` refers to a node with the `end` value on that level that is strictly less than the target range `end` (lines 13–29).

While traversing, it repeatedly snips out marked nodes from the current level as they are encountered (lines 15–22) using a `compareAndSet()`. `compareAndSet()` function also validates that the next field of the predecessor still references the current Node.

Once an unmarked `curr` node is found (line 23), it is tested to see if its `start` is greater than or equal to the target `start`. If so, `pred` is advanced to `curr`, `curr` is advanced to `succ`, and the traverse continues. Otherwise, the current range of `pred` is the immediate predecessor of the target node. The `find()` method then breaks out of the current level search loop, saving the current values of `pred` and `curr` (lines 26–32).

The `find()` method continues this process until it reaches the bottom level. An important point is that each level's traversal maintains the previously described properties. Specifically, if a node with the target key is in the list, it will be found at the bottom level even if nodes are removed at higher levels. When traversal stops, `pred` refers to a predecessor of the target node. The method descends to each next lower level without

skipping over the target node. If the Node is in the list, it will be found at the bottom level. Additionally, if the Node is found, it cannot be marked because if it were marked, it would have been snipped out in lines 15–22. Thus, the condition test on line 35 only needs to check if there are overlap ranges (`findInsert`) or if the start and end of the Node match the target start and end (`findExact`).

1. `findInsert`:

```
return (!(start > pred->getEnd() && end < curr->getStart()));
```

2. `findExact`:

```
return (start == curr->getStart() && end == curr->getEnd());
```

The linearization points for both successful and unsuccessful calls to the `find()` method occur when the `curr` reference at the bottom-level list is set, either at line 11 or line 20, for the last time before the success or failure of the `find()` call is determined at line 35.

```
1  bool ConcurrentRangeLock::releaseLock(uint64_t start, uint64_t end) {
2      Node *preds[maxLevel + 1];
3      Node *succs[maxLevel + 1];
4      Node *succ;
5
6      while (true) {
7          bool found = findExact(start, end, preds, succs);
8          if (!found) {
9              return false;
10         } else {
11             Node *nodeToRemove = succs[0];
12             for (int level = nodeToRemove->getTopLevel();
13                     level >= 0 + 1; level--) {
14                 bool marked[1] = {false};
15                 succ = nodeToRemove->next[level]->get(marked);
16                 while (!marked[0]) {
17                     nodeToRemove->next[level]->attemptMark(succ, true);
18                     succ = nodeToRemove->next[level]->get(marked);
19                 }
20             }
21
22             bool marked[1] = {false};
23             succ = nodeToRemove->next[0]->get(marked);
24             while (true) {
25                 bool iMarkedIt = nodeToRemove->next[0]->compareAndSet(
26                         succ, succ, false, true);
27                 succ = succs[0]->next[0]->get(marked);
28                 if (iMarkedIt) {
29                     findDelete(start, end);
30
31                     return true;
32                 } else if (marked[0]) {
33                     return false;
34                 }
35             }
36         }
37     }
38 }
```

Listing 4.3: Pseudocode for releaseLock method

```
1  bool ConcurrentRangeLock::find(uint64_t start, uint64_t end,
2      Node **preds, Node **succs) {
3      bool marked[1] = {false};
4      bool snip;
5      Node *pred, *succ, *curr;
6
7   retry:
8      while (true) {
9      pred = head;
10     for (int level = maxLevel; level >= 0; level--) {
11         curr = pred->next[level]->getReference();
12
13         while (start > curr->getStart()) {
14             succ = curr->next[level]->get(marked);
15             while (marked[0]) {
16                 snip = pred->next[level]->compareAndSet(curr, succ, false, false);
17
18                 if (!snip) goto retry;
19
20                 curr = pred->next[level]->getReference();
21                 succ = curr->next[level]->get(marked);
22             }
23             if (start >= curr->getStart()) {
24                 pred = curr;
25                 curr = succ;
26             } else {
27                 break;
28             }
29         }
30
31         preds[level] = pred;
32         succs[level] = curr;
33     }
34
35     return **condition**;
36     }
37 }
```

Listing 4.4: General pseudocode for find methods

5 Evaluation

This section will evaluate our proposed concurrent range lock under different scenarios. The goal is to see the scalability and throughput of our mechanism compared to state-of-the-art range locks.

5.1 Competitors

We denote our primary implementation as `Our`. In addition to `Our`, we implement two different competitors for comparison. The first competitor is a scalable range lock proposed by Kogan et al. [2], denoted as `Scalable RL`. We specifically implemented the Exclusive Access variant presented in their paper, as it aligns with the focus of our research. The second competitor is a skip list range lock proposed by Song et al. [3], which we denote as `Song RL`.

5.2 Benchmark enviroment

For benchmark purposes, we utilized a server with an AMD Ryzen 9 7950X processor, which features 16 cores and 32 threads, providing substantial computational power for the experiments. The server configuration included a virtual memory cache with a capacity of 128 GB, backed by 32 GB of physical memory. The cache management aimed to minimize eviction sizes to enhance performance and reliability during the benchmarks.

5.3 Microbenchmark

5.3.1 Workload

The primary objective of a concurrent range lock mechanism is to enable multiple threads to access disjoint parts of the same shared object efficiently. To simulate and evaluate the effectiveness of different range locking strategies, we utilize the `mmap()`

system call to create a shared object in memory and the `memset()` function to simulate write operations to this shared object. For each range, we will write 1KB. The use of `memset()` serves as a placeholder for actual modifications, enabling us to focus on the performance characteristics of the locking mechanism itself.

To explore various levels of contention and potential usage scenarios for range locks, we have devised two distinct workloads, each designed to stress the locking mechanism under different conditions:

- `W1`: In this workload, each thread operates with fine granularity. A thread locks a single memory range, performs a modification (simulated by `memset()`), and immediately releases the lock before proceeding to the next range in its queue. This approach simulates a scenario with minimal contention, focusing on the efficiency of the lock mechanism in handling rapid lock acquisition and release cycles. The primary objective is to assess the overhead introduced by the locking mechanism under low contention and to evaluate its performance in scenarios demanding high throughput with minimal waiting times.

- `W2`: This workload introduces a more complex and realistic scenario. Here, threads perform batched memory operations, where a series of memory ranges (typically 16) are locked, modified (simulated by `memset()`) and then unlocked in a single batch. The goal is to test all competitors under heavier threaded conditions. By locking and unlocking in batches, the number of ranges within the data structure increases, providing insight into how efficiently each competitor can search, lock, and unlock ranges as the load increases. This workload emphasizes the performance of the locking mechanism in handling more substantial, real-world use cases where contention may be higher.

Through these workloads, we aim to comprehensively evaluate the performance of the concurrent range lock mechanism under different workloads and contention scenarios. Each workload provides insights into the lock's scalability, efficiency, and overall robustness in handling varying degrees of parallelism and contention.

5.3.2 Optimal Height for Range Locking

The skip list height in our range-locking mechanism plays a critical role in balancing performance and resource utilization. A skip list with insufficient height may need to optimize lookup times effectively, leading to slower performance. In contrast, an excessive height increases memory consumption and management overhead without proportionate gains in efficiency. This test aims to determine the optimal height for our concurrent range-locking mechanism before comparing it to alternative approaches.

We conducted experiments using the two workloads described in Subsection 5.3.1 to identify the best skip list height.

Runtime on different heights

Figure 5.1: Runtime for different heights of range lock

For **W1**, where contention is minimal, and each thread locks and releases ranges immediately, the height of the skip list has negligible impact on performance. In this scenario, since the number of ranges locked at any given time is limited to the number of active worker threads, the benefits of a taller skip list are not realized. The overhead of managing additional levels outweighs any potential gains, leading to similar performance across different heights.

For **W2**, the total runtime decreases as the skip list height increases, aligning with the expected behavior. The multi-level structure of skip lists allows for faster lookups, reducing runtime as height increases. However, beyond a height of 10, further increases yield diminishing returns. This is due to the probabilistic level assignment in skip lists, where the chance of a node reaching higher levels decreases exponentially (e.g., level 12 has a probability of about $\frac{1}{2^{12}} \approx 0.00024\%$). The sparsely populated upper levels

contribute little to performance improvement but increase memory usage.

Based on these observations, we identify a sweet spot at a height of 10, where the trade-off between performance and resource utilization is optimized. This height balances the need for efficient lookups with manageable memory overhead, making it the ideal choice for our range-locking mechanism in subsequent benchmarks.

5.3.3 Result

Figure 5.2 shows the number of successful locks and unlocks per second with increasing working threads under **W1**. We can see that `Our` has good scalability and outperforms the other two. We archive four times more operations than `Scalable RL` and twelve times more than `Song RL`. The poor performance of `Song RL` is due to the immediate locking and releasing mechanism, combined with the low memory overhead of `memset` (1KB per range). The spinlock in `Song RL` becomes a significant point of contention, leading to its poor performance. Additionally, since the locks are released immediately, `Song RL` cannot effectively leverage its skip list data structure, as the maximum number of ranges in the list is limited to the number of worker threads. However, in more realistic scenarios, such as those discussed in Section 5.4 or **W2**, we will see `Song RL` perform much better.

Figure 5.3 illustrates the same result for **W2**, `Our` continued to outperform the others, achieving two to nine times more operations than `Scalable RL` and two to six times more than `Song RL`. In this workload, the number of ranges locked simultaneously increases to the number of threads multiplied by the batch size (16). This allowed `Song RL` to perform significantly better compared to `Scalable RL`. However, `Scalable RL` demonstrated better scalability as the number of worker threads increased. Beyond sixteen threads, `Scalable RL` began to surpass `Song RL` in performance.

In summary, `Our` consistently outperforms both `Scalable RL` and `Song RL` across different workloads. While `Song RL` struggles with contention in **W1**, it performs better in more complex scenarios like **W2**. However, `Scalable RL` demonstrates better scalability with increasing thread counts, eventually surpassing `Song RL` as threads increase.

Figure 5.2: **Workload 1**: Lock and unlock operations per second


Figure 5.3: **Workload 2**: Lock and unlock operations per second

5.4 Evaluation with real use-cases

While microbenchmarks are useful for evaluating and comparing range locks at algorithm level, they often do not capture the complexities of real-world workloads. Testing in real system allows us to move beyond synthetic benchmarks and evaluate the effectiveness of range locks in a live, operational environment where factors like transaction throughput, latency, and contention can significantly impact performance. This ensures that the range locks are not only theoretically efficient but also practically robust under realistic conditions. For those reason, we run an other benchmark of range lock integrated into `Leanstore`.

5.4.1 Leanstore

`Leanstore` [18] is a high-performance storage engine designed to support various database management systems. To further enhance its capabilities, Nguyen et al. [19] introduced a comprehensive design for allocating and logging large objects, which has been integrated into Leanstore. Their performance study demonstrates that this approach not only outperforms many popular file systems but also ensures transactional consistency and durability for large objects. Given the crucial role that range locks play in their design, we integrated our concurrent range lock mechanism to enable realistic and rigorous benchmarking.

In the proposed design, range locks synchronize access to shared aliasing areas, which are contiguous ranges of virtual memory addresses used to present disjointed extents as contiguous memory. When a worker needs to allocate virtual memory for large BLOBs, particularly when these BLOBs exceed the size of the worker-local aliasing area, it must reserve free virtual memory from the shared aliasing area. A range lock mechanism prevents concurrent workers from accessing overlapping memory regions. The range lock operates by locking specific ranges within the shared aliasing area, ensuring that only one worker can modify or access a particular memory range. This prevents race conditions and ensures data consistency while simultaneously allowing multiple workers to operate on different memory ranges.

5.4.2 Competitors

We integrated all three versions described in Section 5.3. Additionally, we included the original, specifically optimized range lock version by the author as a fourth competitor.

The range lock manages locking over a range of blocks in a buffer. It iterates through the specified block range in chunks, calculating which bits in the lock need to be updated for each chunk. It tries to set the corresponding bits case lock and clear bits case unlock in an atomic manner. If a bit was already set by another thread, the function detects this conflict, aborts the current operation, and undoes any changes made so far.

This range lock is lightweight and efficient, especially for managing small to moderate numbers of logical blocks in memory. However, it depends on specific memory size constraints. Since our research focuses exclusively on the generalization use case of range lock, we have chosen not to consider this approach as a competitor in our project.

5.4.3 Workload

The experiment, utilizing synthetic YCSB [20] workloads, was designed to evaluate the performance of LeanStore under a read-intensive, multithreaded environment. The workload comprises exclusively read operations, executed over a 10-second duration, with each read operation employing a straightforward `memcpy()` function. Each record in the workload has a payload size of 1 MB, with 1000 records being processed. The experiment also incorporates a buffer management configuration that allocates 128 GB of virtual and 32 GB of physical memory, providing a robust test of LeanStore's ability to manage large data sets under constrained physical memory conditions.

5.4.4 Result

We evaluated the throughput of four competitors as the number of worker threads increased, as depicted in Figure 5.4. As anticipated, the `Bitmap` implementation consistently outperformed the other three competitors. This superiority can be attributed mainly to its optimized use of the aliasing arena, making it the most efficient workload handling. Our implementation (`Our`) demonstrated competitive performance, closely trailing `Bitmap` in terms of transactions per second. Both `Bitmap` and `Our` showed effective scalability with an increasing number of worker threads, though `Bitmap` maintained a consistent, albeit slight, performance advantage.

Conversely, `Song RL` initially exhibited substantial throughput but encountered a sharp decline beyond 24 worker threads. This decline underscores the limitations of its coarse-grained range lock mechanism, as discussed in Subsection 2.1.2. The single spinlock in `Song RL` becomes a significant bottleneck in high-contention scenarios, leading to marked performance degradation. On the other hand, the `Scalable RL` showed steady,

Figure 5.4: Leanstore YCSB: Transactional Throughput [Txn/s].

linear throughput growth. However, its overall performance remained considerably lower than the other competitors, highlighting its limitations.

An additional key observation is the difficulty faced by the `Bitmap`, `Our`, and `Song RL` competitors in scaling effectively beyond 16 worker threads. This limitation is due mainly to the BLOB size in our workload. With a BLOB size of 1MB and 16 worker threads, the combined size of the client-side buffer and the internal DBMS memory block for the BLOB exceeds the L3 cache capacity (32MB in our machine), resulting in significant contention at the L3 cache. Furthermore, LeanStore's use of `memcpy` for read operations exacerbates this issue. `memcpy` consumes 2MB of memory read and write for every 1MB BLOB. An average throughput of approximately 60,000 operations per second equals over 110 GB of memory consumed by `memcpy`, which saturates the memory bandwidth and hinders the application's ability to scale.

We also analyzed other performance metrics, such as cycles, instructions, and L1 cache misses. As illustrated in Figure 5.5, both the `Bitmap` and `Our` implementations maintained stable levels of cycles, instructions, and L1 cache misses as the number

of threads increased. In contrast, `Song RL` exhibited a significant rise in the number of instructions with additional threads, indicating inefficiencies under higher loads. Meanwhile, `Scalable RL` demonstrated better scalability, with a decreasing number of L1 misses and instructions as the workload intensified.

(a) Cycles by thread workers

(b) Instructions by thread workers

(c) L1-misses by thread workers

Figure 5.5: Cycles, instructions, and L1-misses

6 Conclusion

This thesis presented a new lock-free concurrent range lock to address the bottlenecks and inefficiencies in existing range-locking methods. Traditional techniques, such as coarse-grained locking and tree or list-based locks, often struggle in high-performance computing environments due to contention and poor scalability under heavy concurrency. We introduced a new approach based on a probabilistic skip list to tackle these challenges. We utilized atomic operations like `compareAndSet` to manage node references without relying on traditional locking mechanisms. This design significantly reduces bottlenecks and enhances concurrency performance, making it more suitable for large-scale distributed systems.

Our extensive evaluation, including microbenchmarks and integration into the Leanstore storage engine, demonstrated the superiority of our approach over existing range-locking methods. The results highlighted notable improvements in throughput and reduced contention across various workloads and thread counts, affirming the effectiveness of our mechanism in real-world, high-concurrency scenarios.

Despite these advancements, there are areas that warrant further exploration and optimization. One of the key challenges identified during our benchmarking is that our implementation, while effective, needs to scale better when dealing with workloads beyond L3 cache lines. This limitation suggests a need for further enhancement of our approach, particularly in terms of cache efficiency, to ensure better scalability in high-performance environments.

Another critical area for future work is the development of more efficient memory deallocation strategies. The current design lacks a robust mechanism for managing memory deallocation, which could lead to potential memory leaks or inefficient memory usage over time. Addressing this issue is crucial to improving the overall efficiency and reliability of our range-locking mechanism.

Additionally, we plan to conduct a detailed analysis of the memory overhead associated with our implementation. The use of additional layers of linked lists and abstractions like `AtomicMarkableReference` introduces a certain level of memory overhead. By thoroughly analyzing this aspect and comparing it with other competitors, we aim to

gain a better understanding of the trade-offs between memory usage and performance, which will help us refine our approach further.

Furthermore, we want to explore how our concurrent range lock can be integrated with other data structures or applied in different contexts. For example, investigating the potential benefits of incorporating techniques used in the lock tree implementation in PerconaFT [21] could lead to new opportunities for enhancing the flexibility and applicability of our approach. Exploring these avenues could open up new possibilities for improving our design and expanding its utility across various high-concurrency environments.

Bibliography

[1] J. Gao, Y. Lu, M. Xie, Q. Wang, and J. Shu. "Citron: Distributed Range Lock Management with One-sided {RDMA}". In: *21st USENIX Conference on File and Storage Technologies (FAST 23)*. 2023, pp. 297–314.

[2] A. Kogan, D. Dice, and S. Issa. "Scalable range locks for scalable address spaces and beyond". In: *Proceedings of the Fifteenth European Conference on Computer Systems*. 2020, pp. 1–15.

[3] X. Song, J. Shi, R. Liu, J. Yang, and H. Chen. "Parallelizing live migration of virtual machines". In: *Proceedings of the 9th ACM SIGPLAN/SIGOPS international conference on Virtual execution environments*. 2013, pp. 85–96.

[4] D. Lomet and M. F. Mokbel. "Locking key ranges with unbundled transaction services". In: *VLDB*. 2009.

[5] G. Congiu, S. Narasimhamurthy, T. Süß, and A. Brinkmann. "Improving collective I/O performance using non-volatile memory devices". In: *2016 IEEE International Conference on Cluster Computing (CLUSTER)*. IEEE. 2016, pp. 120–129.

[6] Q. Kang, S. Breitenfeld, K. Hou, W.-k. Liao, R. Ross, and S. Byna. "Optimizing performance of parallel I/O accesses to non-contiguous blocks in multiple array variables". In: *2021 IEEE International Conference on Big Data (Big Data)*. IEEE. 2021, pp. 98–108.

[7] J. Corbet. "Range reader/writer locks for the kernel". In: *LWN.net* (2022). Accessed: 2024-04-21. URL: https://lwn.net/Articles/724502/.

[8] M. Rybczynska. "Introducing maple trees". In: *LWN.net* (2022). Accessed: 2024-04-21. URL: https://lwn.net/Articles/845507/.

[9] J. Corbet. "The ongoing search for mmap_lock scalability". In: *LWN.net* (2022). Accessed: 2024-04-21. URL: https://lwn.net/Articles/893906/.

[10] J. Kara. "Implement range locks". In: *lkml.org* (2013). Accessed: 2024-04-21. URL: https://lkml.org/lkml/2013/1/31/483.

[11] G. Graefe. "Hierarchical locking in B-tree indexes". In: *On Transactional Concurrency Control*. Springer, 2007, pp. 45–73.

[12] A. Pavlo. "Two-Phase Locking". In: *15445.courses.cs.cmu.edu* (2022). Accessed: 2024-04-21. URL: `https://15445.courses.cs.cmu.edu/fall2022/slides/16-twophaselocking.pdf`.

[13] J.-H. Kim, J. Kim, H. Kang, C.-G. Lee, S. Park, and Y. Kim. "pNOVA: Optimizing shared file I/O operations of NVM file system on manycore servers". In: *Proceedings of the 10th ACM SIGOPS Asia-Pacific Workshop on Systems*. 2019, pp. 1–7.

[14] W. Pugh. *A skip list cookbook*. Citeseer, 1990.

[15] M. Fomitchev and E. Ruppert. "Lock-free linked lists and skip lists". In: *Proceedings of the twenty-third annual ACM symposium on Principles of distributed computing*. 2004, pp. 50–59.

[16] W. Pugh. "Skip lists: a probabilistic alternative to balanced trees". In: *Communications of the ACM* 33.6 (1990), pp. 668–676.

[17] H. Maurice, S. Nir, L. Victor, and S. Michael. *The art of multiprocessor programming*. Newnes, 2020.

[18] V. Leis, M. Haubenschild, A. Kemper, and T. Neumann. "LeanStore: In-memory data management beyond main memory". In: *2018 IEEE 34th International Conference on Data Engineering (ICDE)*. IEEE. 2018, pp. 185–196.

[19] L.-D. Nguyen and V. Leis. "Why Files If You Have a DBMS". In: ICDE. 2024.

[20] B. F. Cooper, A. Silberstein, E. Tam, R. Ramakrishnan, and R. Sears. "Benchmarking cloud serving systems with YCSB". In: *Proceedings of the 1st ACM symposium on Cloud computing*. 2010, pp. 143–154.

[21] *PerconaFT*. `https://github.com/percona/PerconaFT`. Accessed: 2024-08-27. 2024.